



Memory and Address Space

Memory Virtualization



- What is memory virtualization?
 - OS virtualizes its physical memory
 - OS provides an illusion memory space per each process
 - It seems to be seen like each process uses the whole memory


Benefit of Memory Virtualization



- Ease of use in programming
- Memory efficiency in terms of times and space
- The guarantee of isolation for processes as well as OS
 - Protection from errant accesses of other processes


OS in The Early System

- Load only one process in memory
 - Poor utilization and efficiency




Multiprogramming and Time Sharing

- **Load multiple processes** in memory
 - Execute one for a short while
 - Switch processes between them in memory
 - Increase utilization and efficiency
- Cause an important **protection issue**
 - Errant memory accesses from other processes




Address Space

- OS creates an abstraction of physical memory.
 - The address space contains all about a running process.
 - That is consist of program code, heap, stack and etc.



Address Space

- Code
 - Where instructions live
- Heap
 - Dynamically allocate memory.
 - malloc in C language
 - new in object-oriented language
- Stack
 - Store return addresses or values.
 - Contain local variables arguments to routines.



Address Space

Virtual Address

- **Every address** in a running program is virtual
 - OS translates the virtual address to physical address

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {

    printf("location of code   : %p\n", (void *) main);
    printf("location of heap   : %p\n", (void *) malloc(1));
    int x = 3;
    printf("location of stack  : %p\n", (void *) &x);

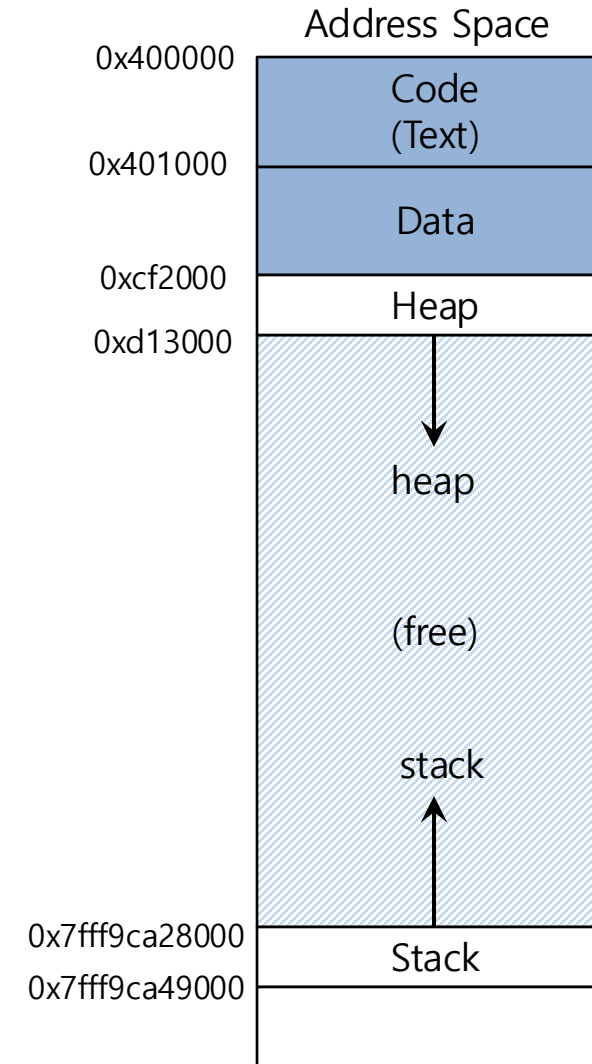
    return x;
}
```

A simple program that prints out addresses

Virtual Address

- The output in 64-bit Linux machine

```
location of code   : 0x40057d  
location of heap   : 0xcf2010  
location of stack  : 0x7fff9ca45fcc
```



Memory API: *malloc()*

```
#include <stdlib.h>

void* malloc(size_t size)
```

- Allocate a memory region on the heap
 - Argument
 - `size_t size` : size of the memory block(in bytes)
 - `size_t` is an unsigned integer type.
 - Return
 - Success : a void type pointer to the memory block allocated by `malloc`
 - Fail : a null pointer

Memory API: *sizeof()*

- Routines and macros are utilized for `size` in `malloc` instead typing in a number directly
- Two types of results of `sizeof` with variables
 - The actual size of '`x`' is known at run-time

```
int *x = malloc(10 * sizeof(int));  
printf("%d\n", sizeof(x));
```

4

- The actual size of '`x`' is known at compile-time

```
int x[10];  
printf("%d\n", sizeof(x));
```

40

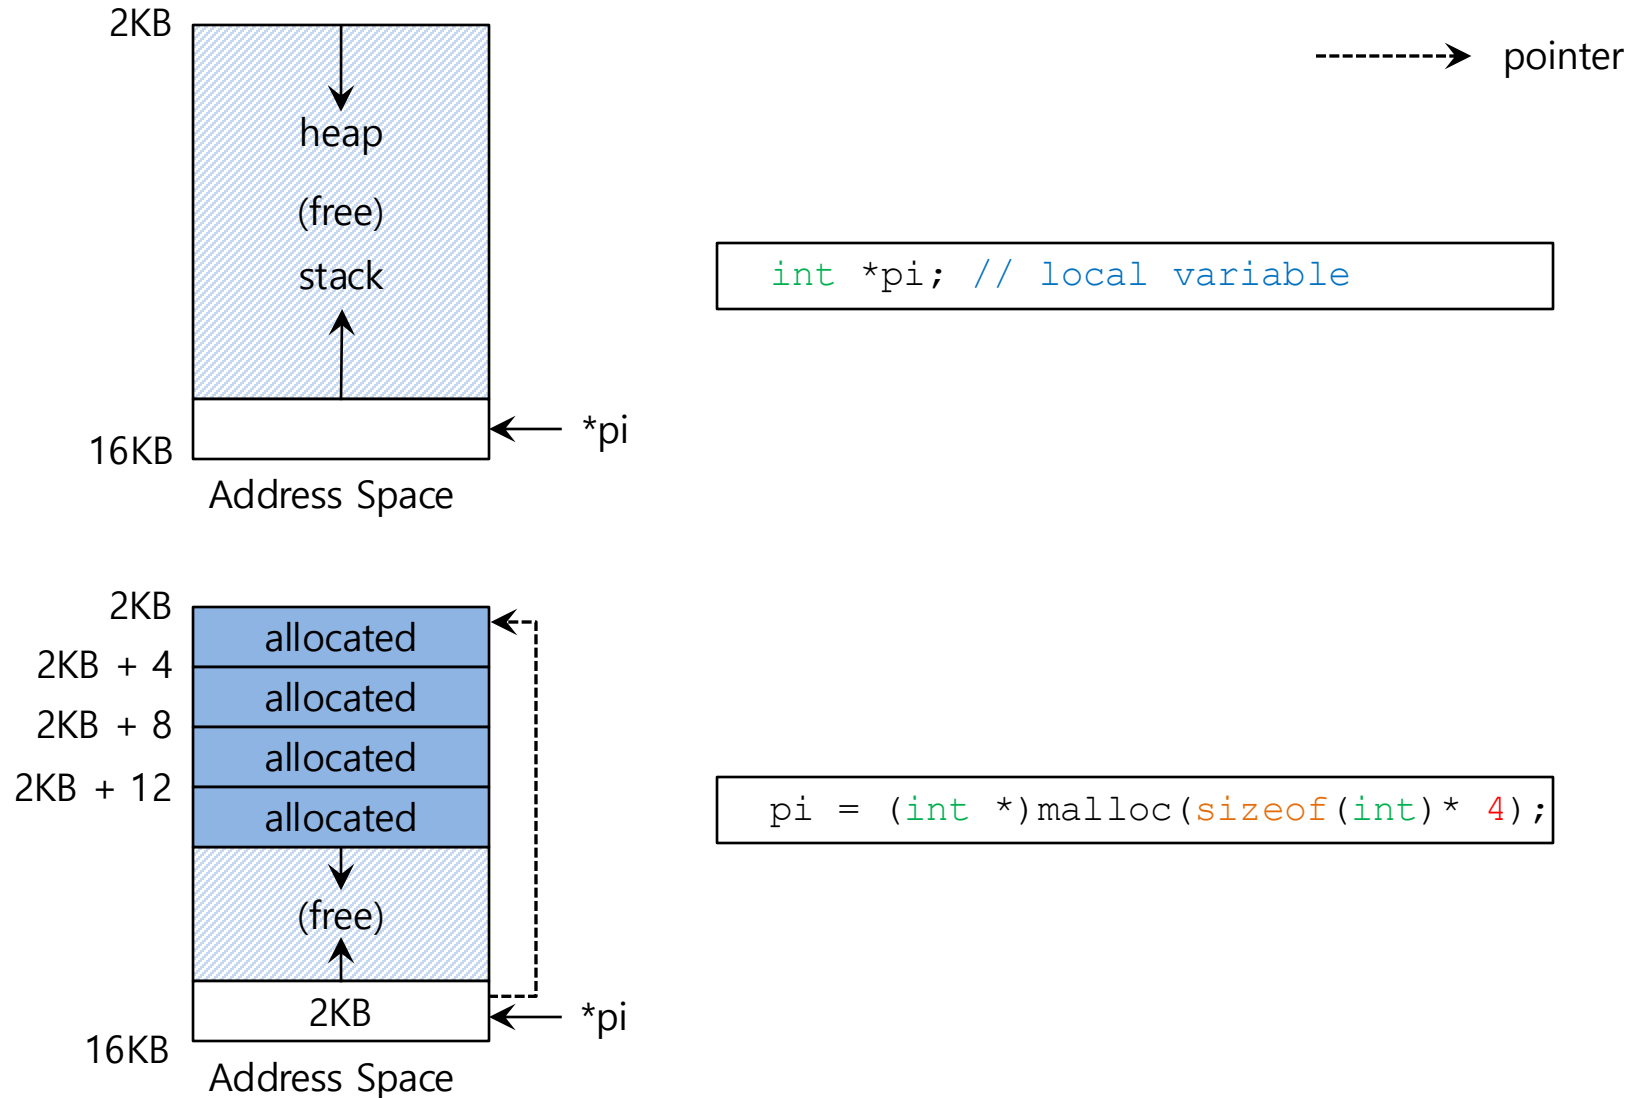
Memory API: *free()*

```
#include <stdlib.h>


void free(void* ptr)
```

- Free a memory region allocated by a call to `malloc`
 - Argument
 - `void *ptr`: a pointer to a memory block allocated with `malloc`
 - Return
 - none


Memory Allocating



Memory Freeing



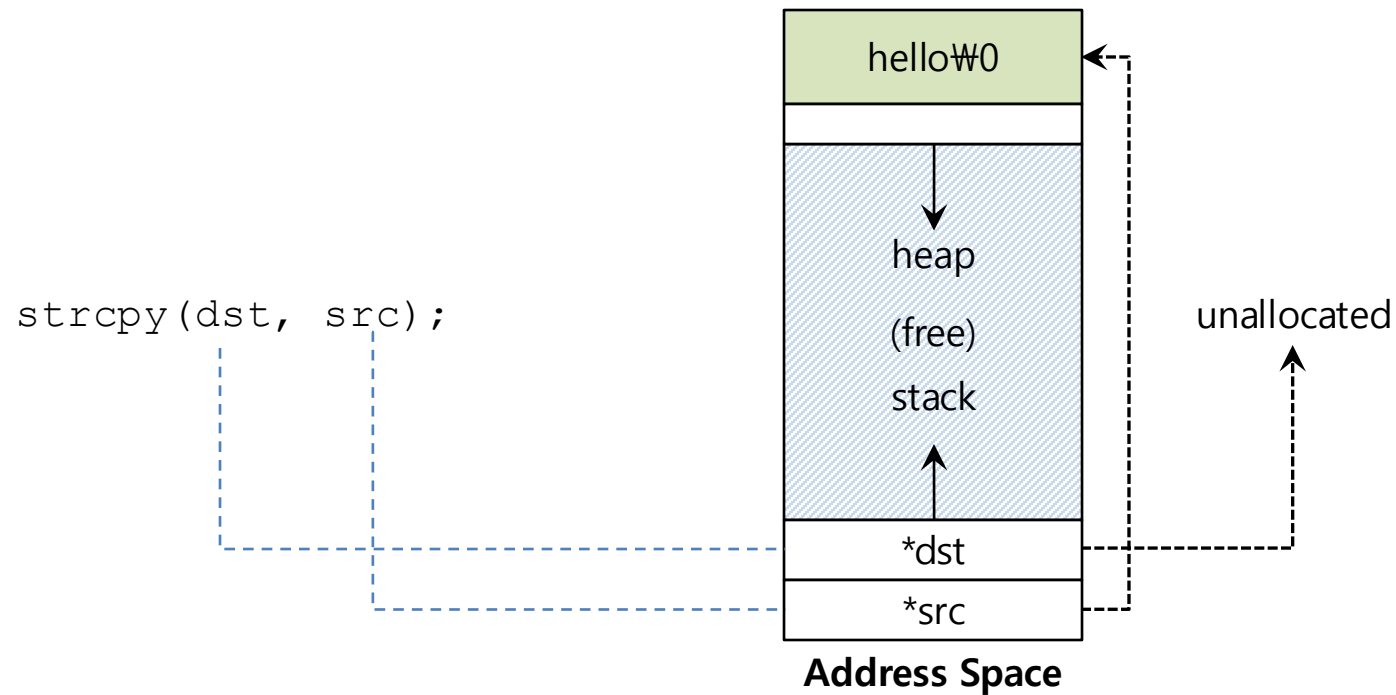
```
free(pi);
```



Forgetting To Allocate Memory

- Incorrect code

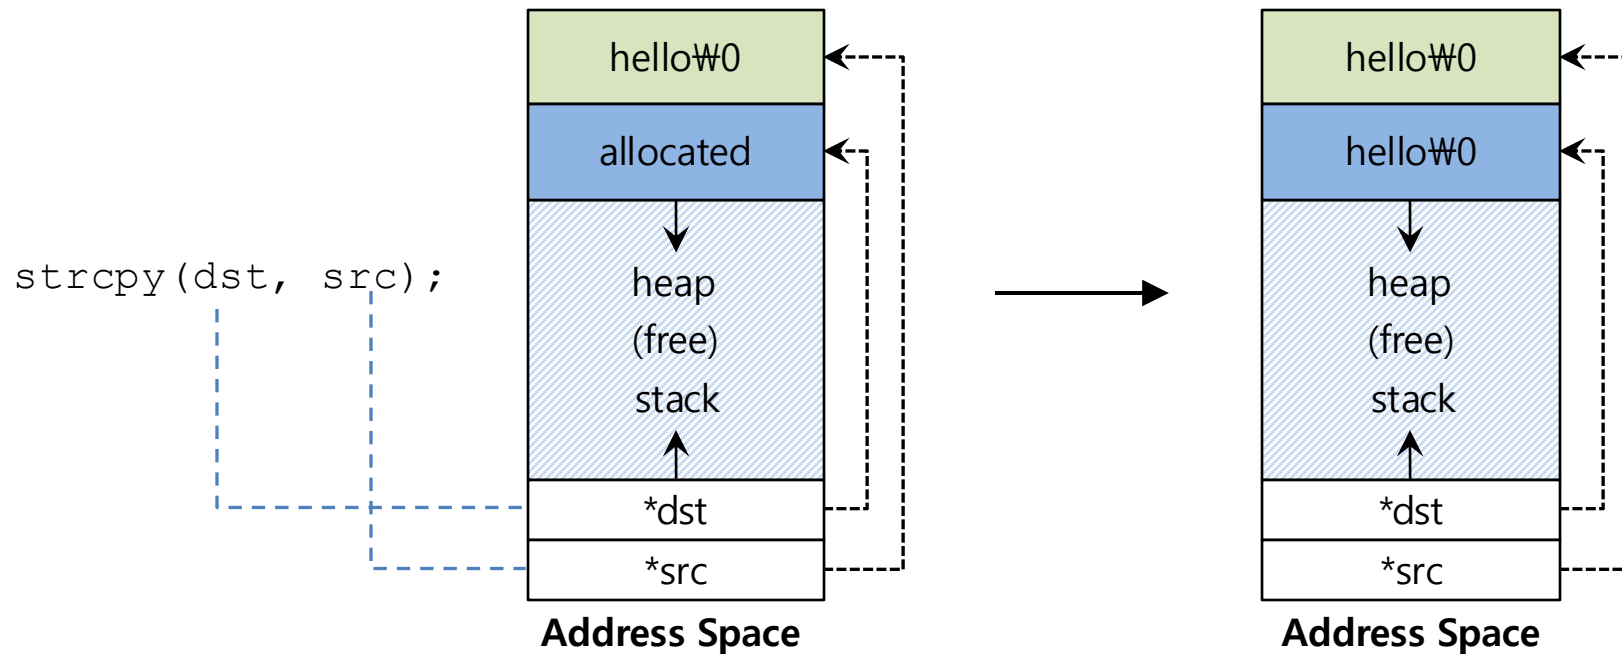
```
char *src = "hello"; //character string constant
char *dst;           //unallocated
strcpy(dst, src);    //segfault and die
```



Forgetting To Allocate Memory

- Correct code

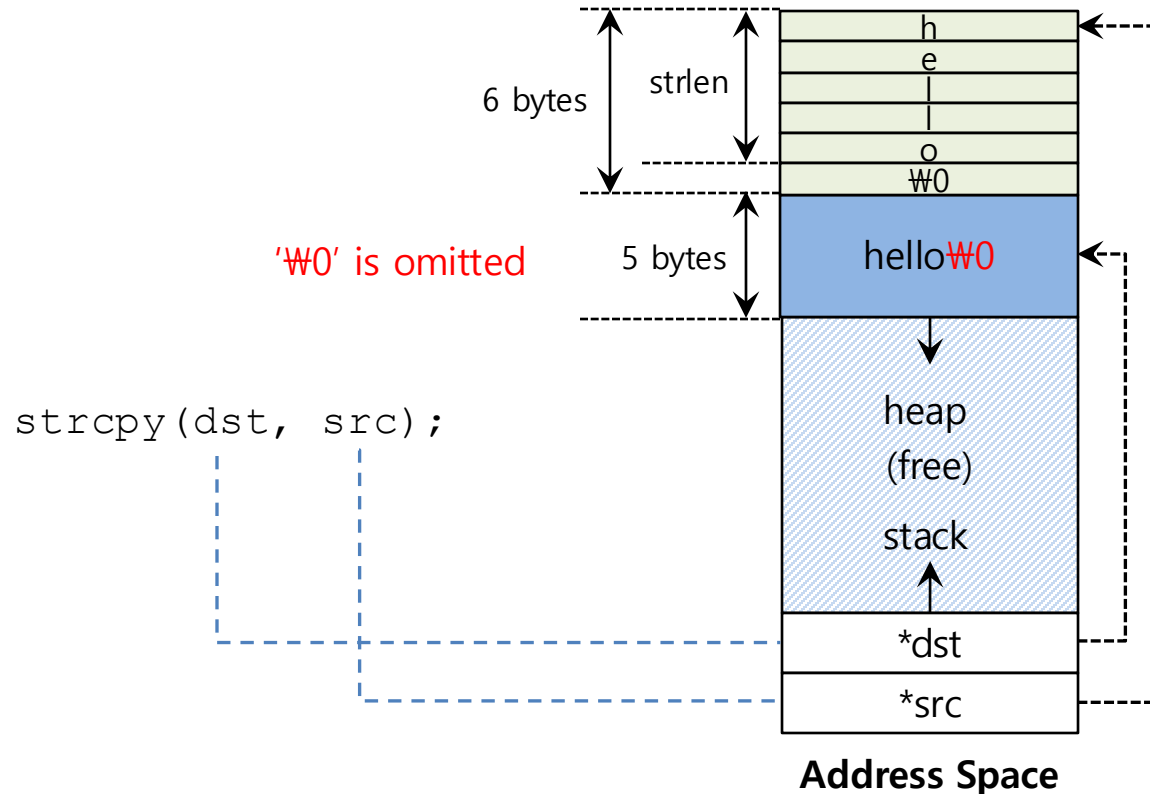
```
char *src = "hello"; //character string constant
char *dst (char *)malloc(strlen(src) + 1 ); // allocated
strcpy(dst, src); //work properly
```



Not Allocating Enough Memory

- Incorrect code, but work properly


```
char *src = "hello"; //character string constant
char *dst (char *)malloc(strlen(src)); // too small
strcpy(dst, src);    //work properly
```




Forgetting to Initialize

- Encounter an uninitialized read

```
int *x = (int *)malloc(sizeof(int)); // allocated
printf("*x = %d\n", *x); // uninitialized memory access
```



Address Space

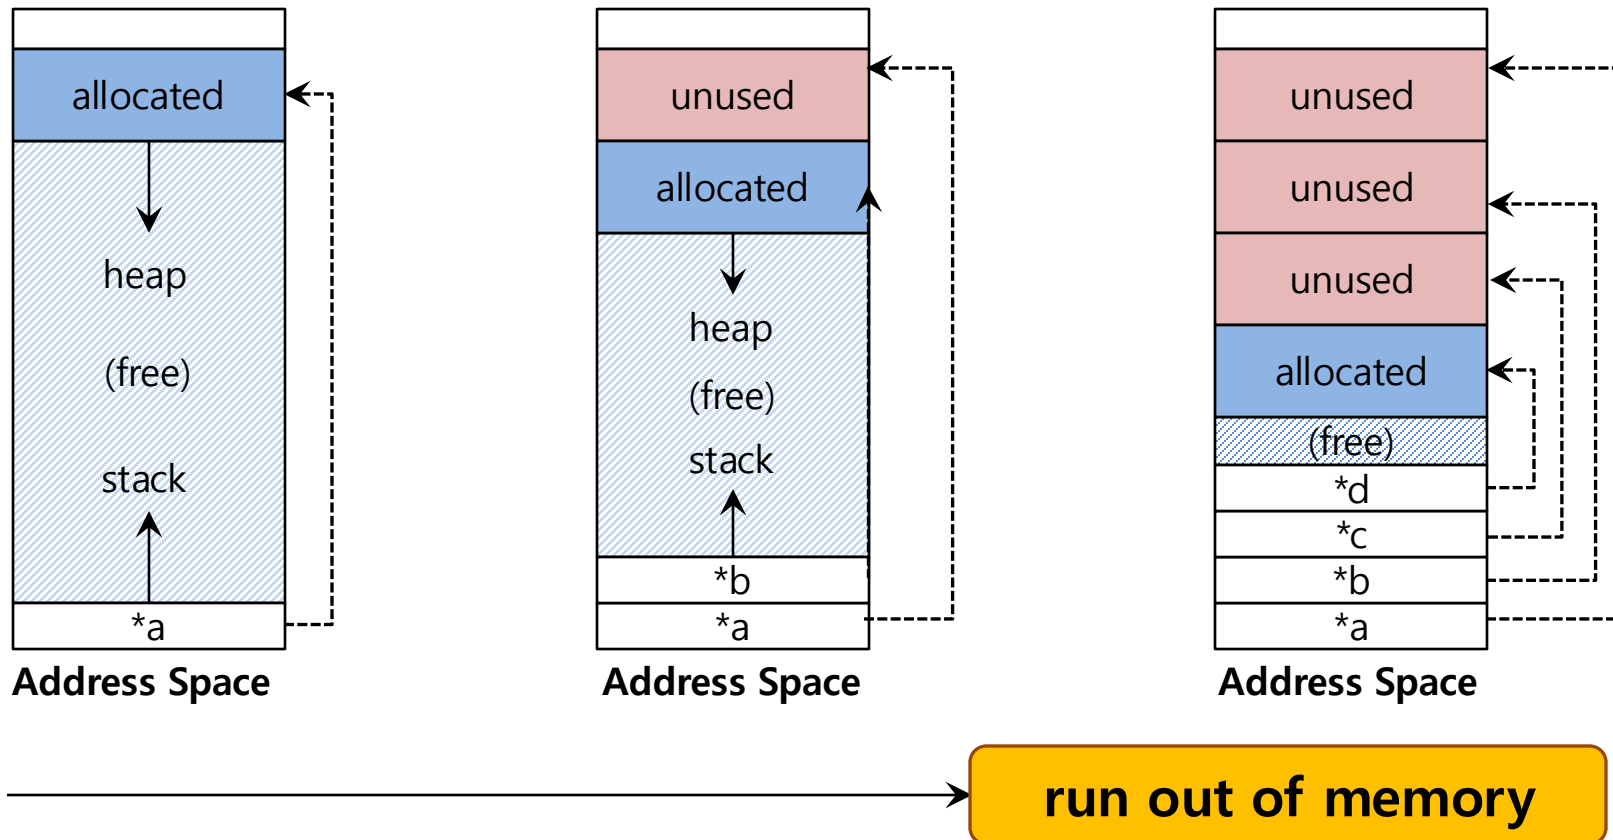


Address Space

Memory Leak




- A program runs out of memory and eventually dies

unused : unused, but not freed




Dangling Pointer

- Freeing memory before it is finished using
 - A program accesses to memory with an invalid pointer



`free(b)`



Other Memory APIs: calloc()

```
#include <stdlib.h>


void *calloc(size_t num, size_t size)
```

- Allocate memory on the heap and zeroes it before returning
 - Argument
 - size_t num : number of blocks to allocate
 - size_t size : size of each block(in bytes)
 - Return
 - Success : a void type pointer to the memory block allocated by calloc
 - Fail : a null pointer

Double Free

- Free memory that was freed already

```
int *x = (int *)malloc(sizeof(int)); // allocated
free(x); // free memory
free(x); // free repeatedly
```



Other Memory APIs: realloc()

```
#include <stdlib.h>

void *realloc(void *ptr, size_t size)
```

- Change the size of memory block

- A pointer returned by `realloc` may be either the same as `ptr` or a new
- Argument
 - `void *ptr`: Pointer to memory block allocated with `malloc`, `calloc` or `realloc`
 - `size_t size`: New size for the memory block(in bytes)
- Return
 - Success: Void type pointer to the memory block
 - Fail : Null pointer

System Calls for Allocation

```
#include <unistd.h>

int brk(void *addr)
void *sbrk(intptr_t increment);
```

- `malloc` library call use `brk` system call
 - `brk` is called to expand the program's *break*
 - *break*: The location of **the end of the heap** in address space
 - `sbrk` is an additional call similar with `brk`
 - Programmers **should never directly call** either `brk` or `sbrk`

System Calls for Memory Mapping

```
#include <sys/mman.h>

void *mmap(void *ptr, size_t length, int port, int flags,
int fd, off_t offset)
```

- `mmap` system call can create **an anonymous** memory region

Memory Virtualizing with Efficiency and Control

- Memory virtualizing takes a similar strategy known as **limited direct execution(LDE)** for efficiency and control
- In memory virtualizing, efficiency and control are attained by **hardware support**
 - e.g., registers, TLB(Translation Look-aside Buffer)s, page-table

Address Translation



- Hardware transforms a **virtual address** to a **physical address**
 - The desired information is actually stored in a physical address
- The OS must get involved at key points to set up the hardware
 - The OS must manage memory to judiciously intervene

Example: Address Translation

- C - Language code

```
void func()  
    int x;  
    ...  
    x = x + 3; // this is the line of code we are interested in
```

- **Load** a value from memory
- **Increment** it by three
- **Store** the value back into memory


Example: Address Translation

■ Assembly

```
128 : movl 0x0(%ebx), %eax      ; load 0+ebx into eax
132 : addl $0x03, %eax          ; add 3 to eax register
135 : movl %eax, 0x0(%ebx)      ; store eax back to mem
```

- Presume that the address of 'x' has been place in `ebx` register.
- **Load** the value at that address into `eax` register.
- **Add** 3 to `eax` register.
- **Store** the value in `eax` back into memory.

Example: Address Translation

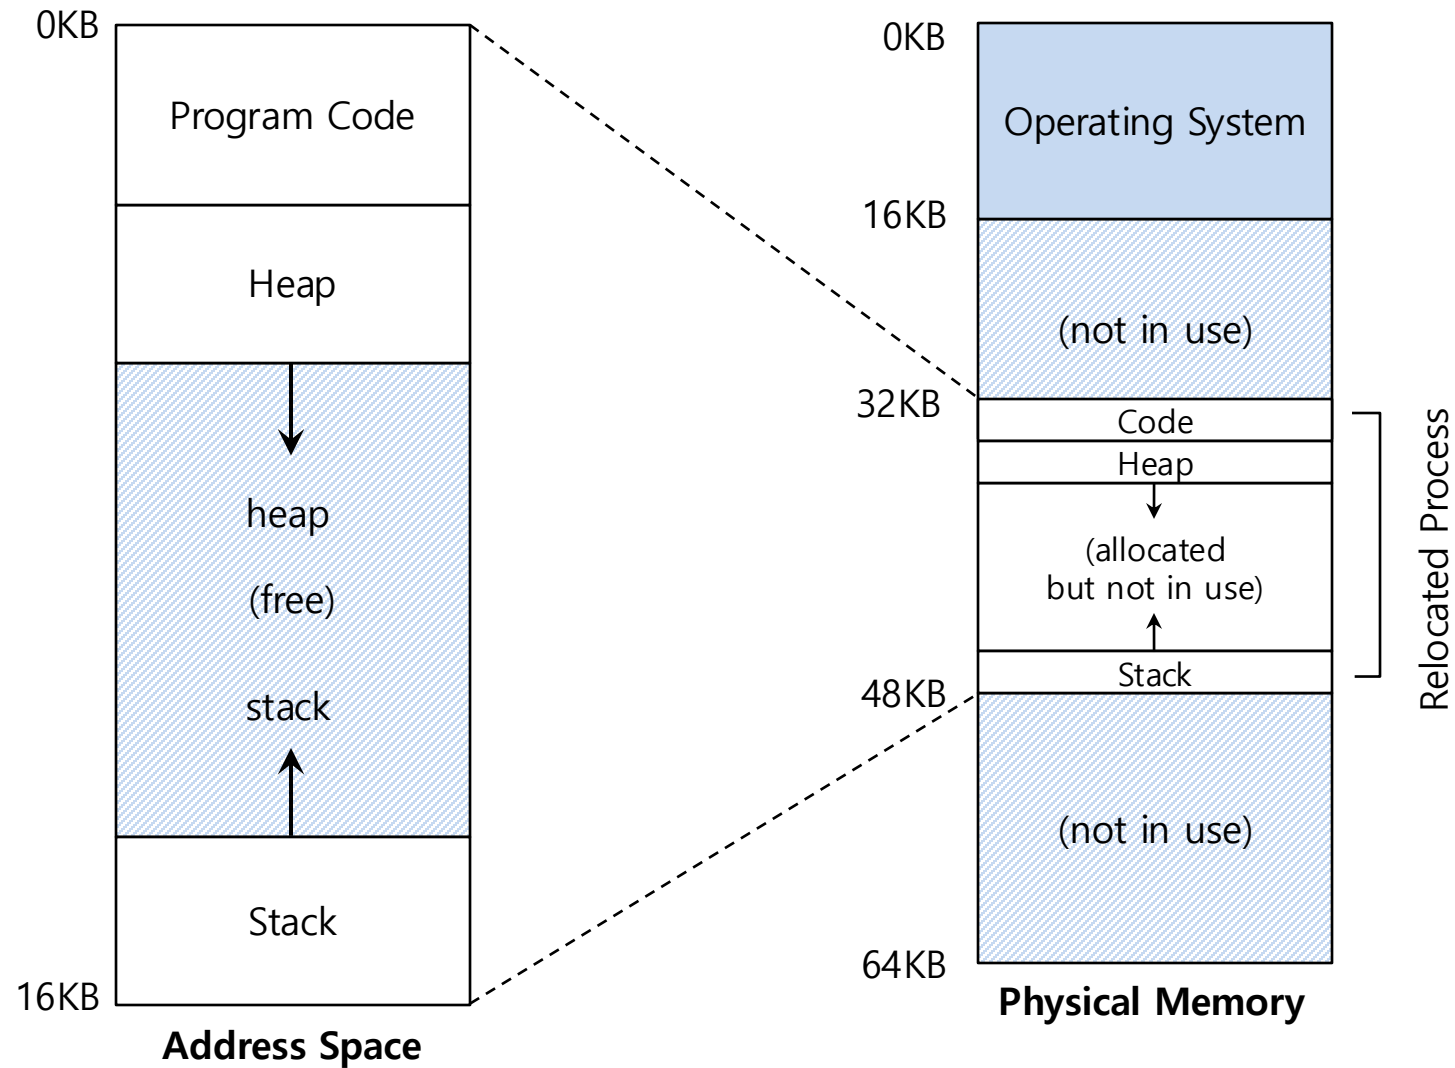


- Fetch instruction at address 128
- Execute this instruction (load from address 15KB)
- Fetch instruction at address 132
- Execute this instruction (no memory reference)
- Fetch the instruction at address 135
- Execute this instruction (store to address 15 KB)

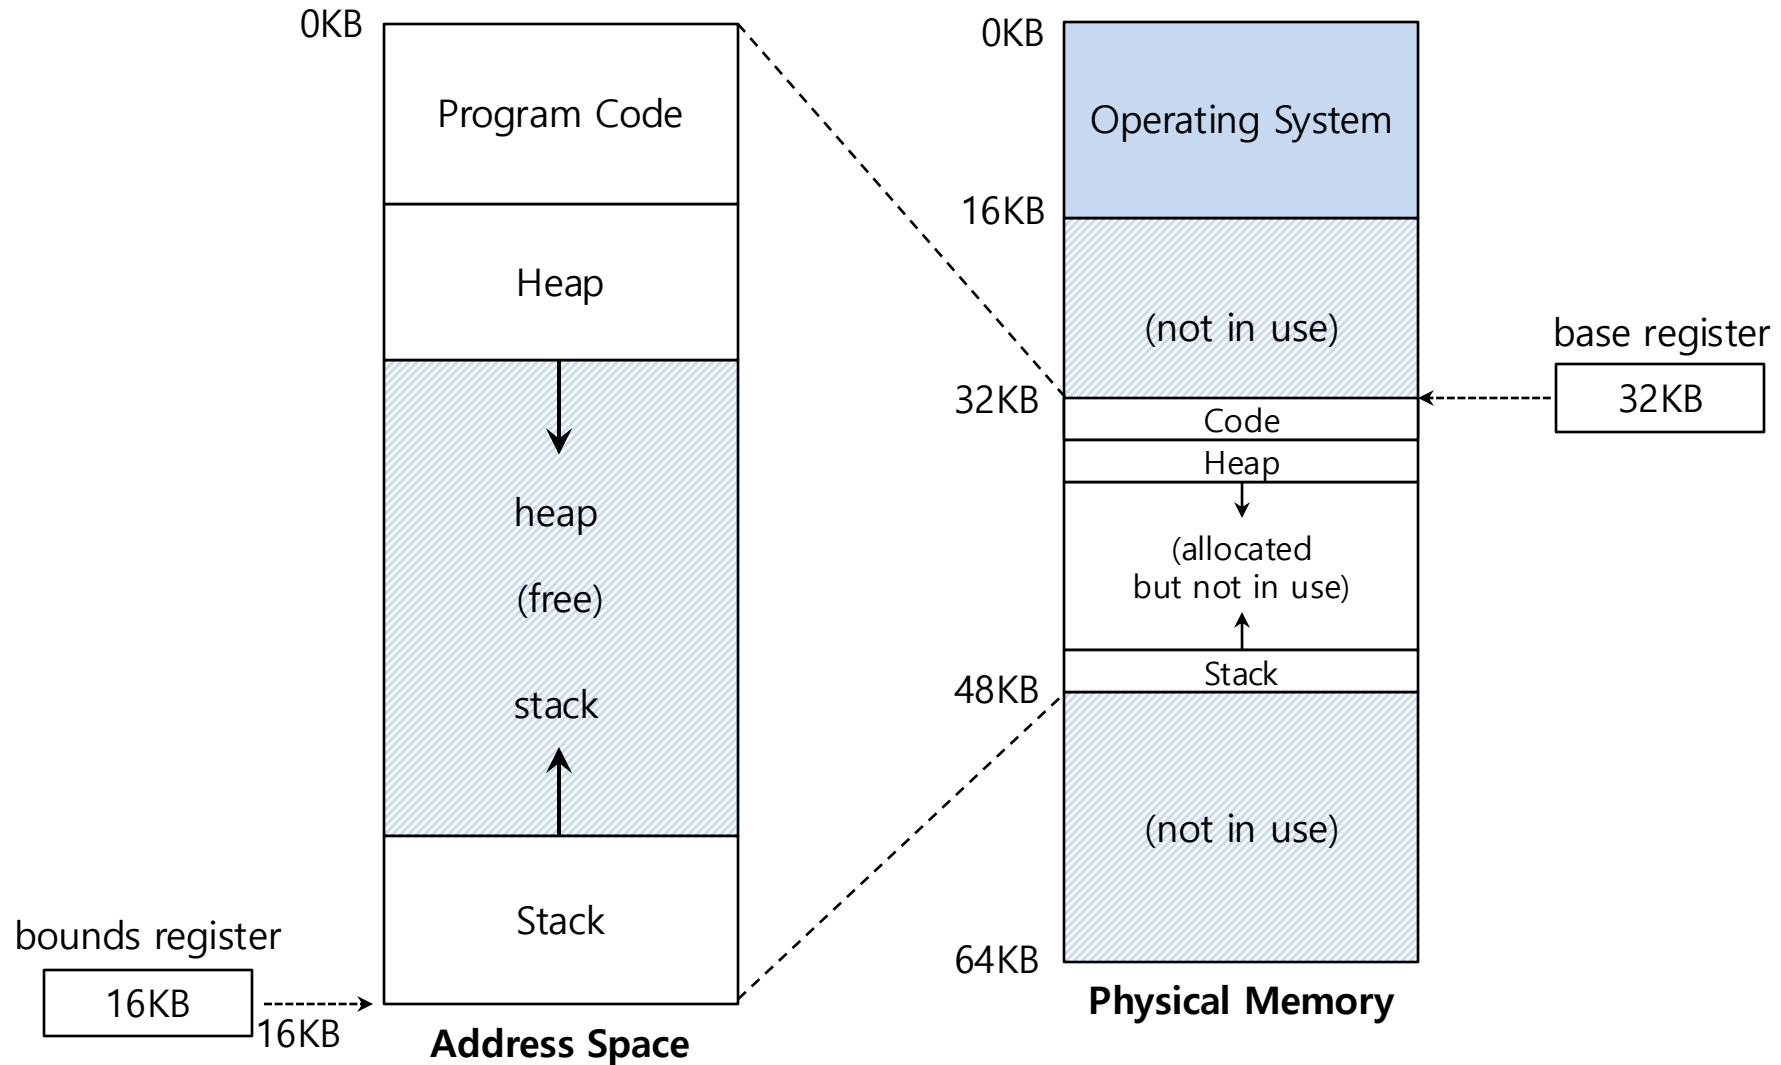
Relocation Address Space

- The OS wants to place the process **somewhere else** in physical memory, not at address 0
 - The address space start at address 0

A Single Relocated Process



Base and Bounds Register



Dynamic(Hardware base) Relocation

- When a program starts running, the OS decides **where** in physical memory a process should be **loaded**
 - Set the **base** register a value

$$\text{physical address} = \text{virtual address} + \text{base}$$

- Every virtual address must **not be greater than bound** and **negative**

$$0 \leq \text{virtual address} < \text{bounds}$$

Relocation and Address Translation


128 : `movl 0x0(%ebx), %eax`

- **Fetch** instruction at address 128

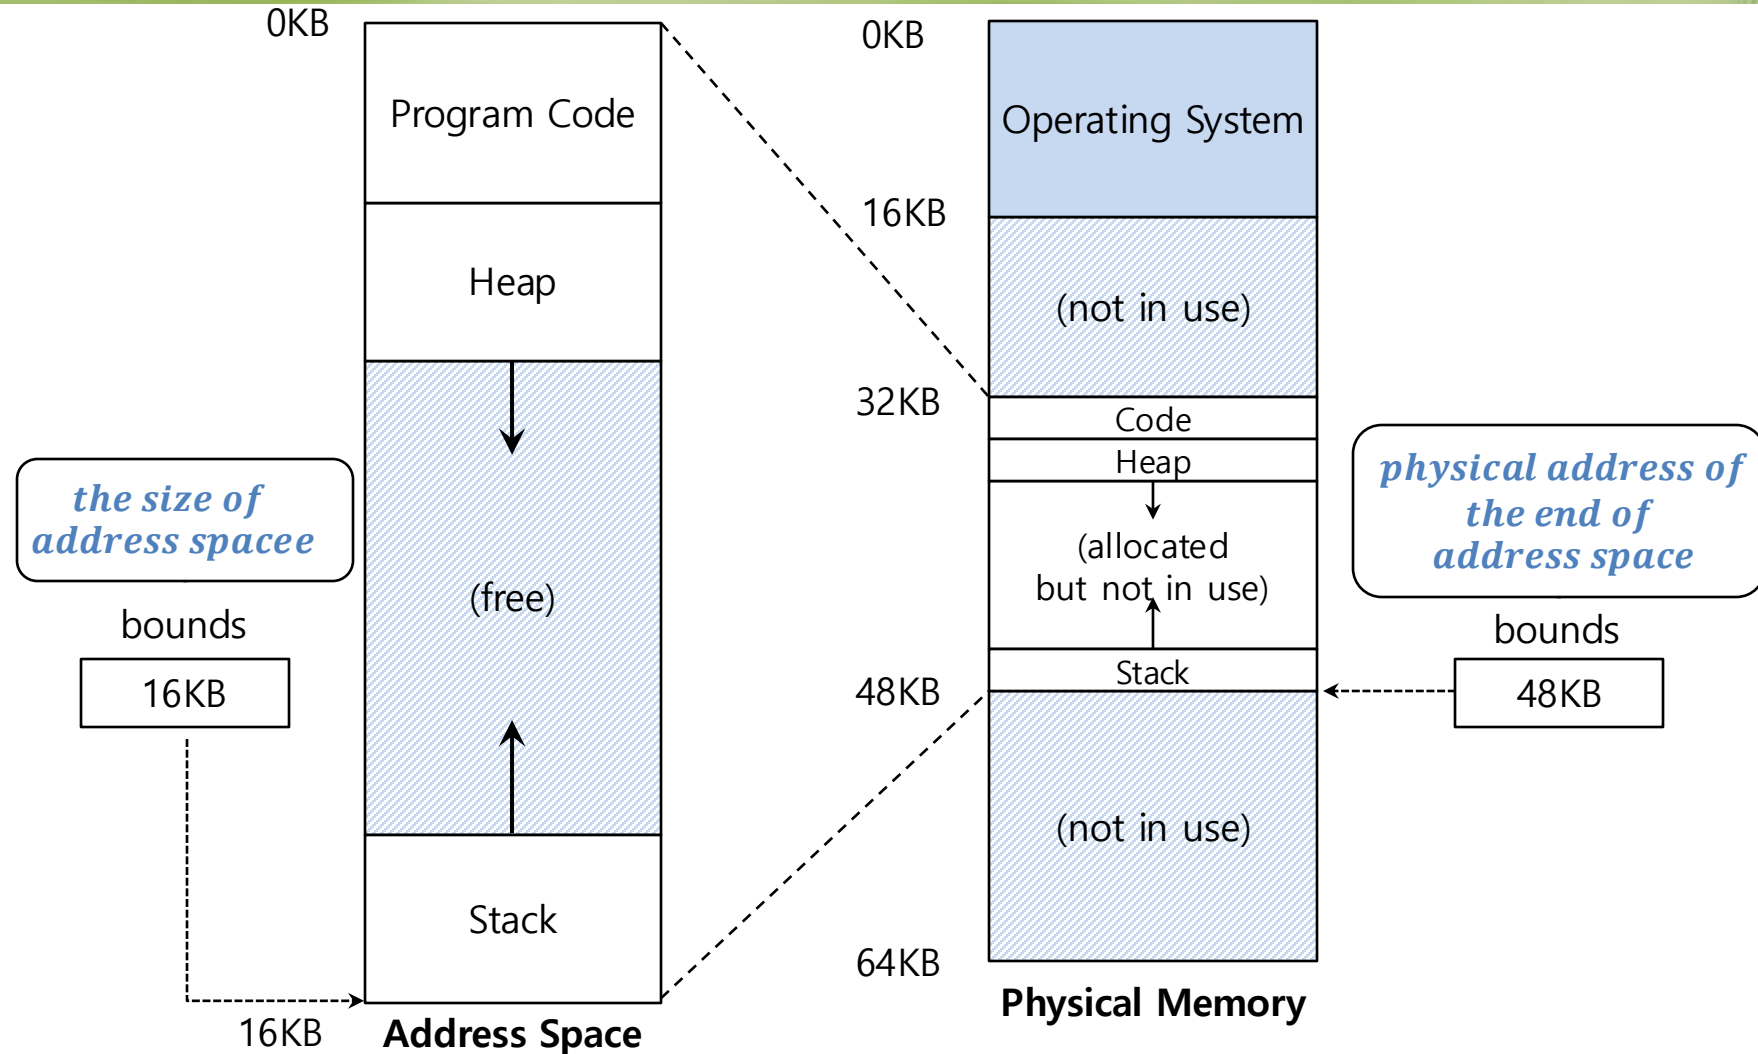
$$32896 = 128 + 32KB(base)$$

- **Execute** this instruction
 - Load from address 15KB

$$47KB = 15KB + 32KB(base)$$



Two ways of Bounds Register



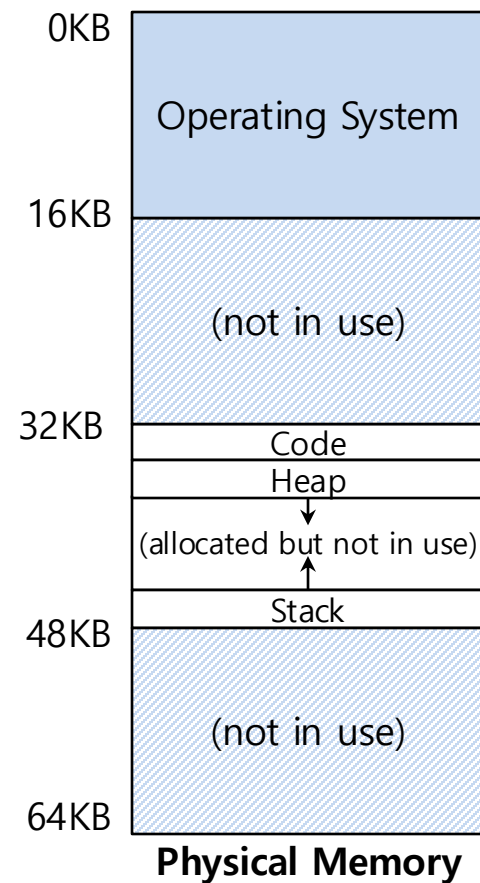

OS Issues for Memory Virtualizing

- The OS must **take action** to implement **base-and-bounds** approach
- Three critical junctures:
 - When a process **starts running**:
 - Finding space for address space in physical memory
 - When a process is **terminated**:
 - Reclaiming the memory for use
 - When context **switch occurs**:
 - Saving and storing the base-and-bounds pair

OS Issues: When a Process Starts Running

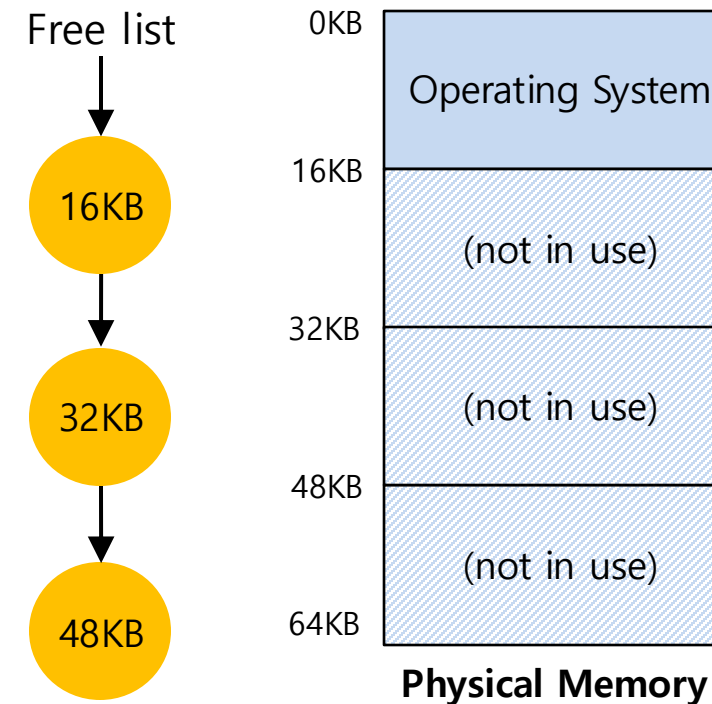

- The OS must **find a room** for a new address space
 - free list : A list of the range of the physical memory which are not in use

The OS lookup the free list



OS Issues: When a Process Is Terminated

- The OS must **put the memory back** on the free list



OS Issues: When Context Switch Occurs

- OS must **save and restore** the base-and-bounds pair
 - In **process structure** or **process control block(PCB)**

