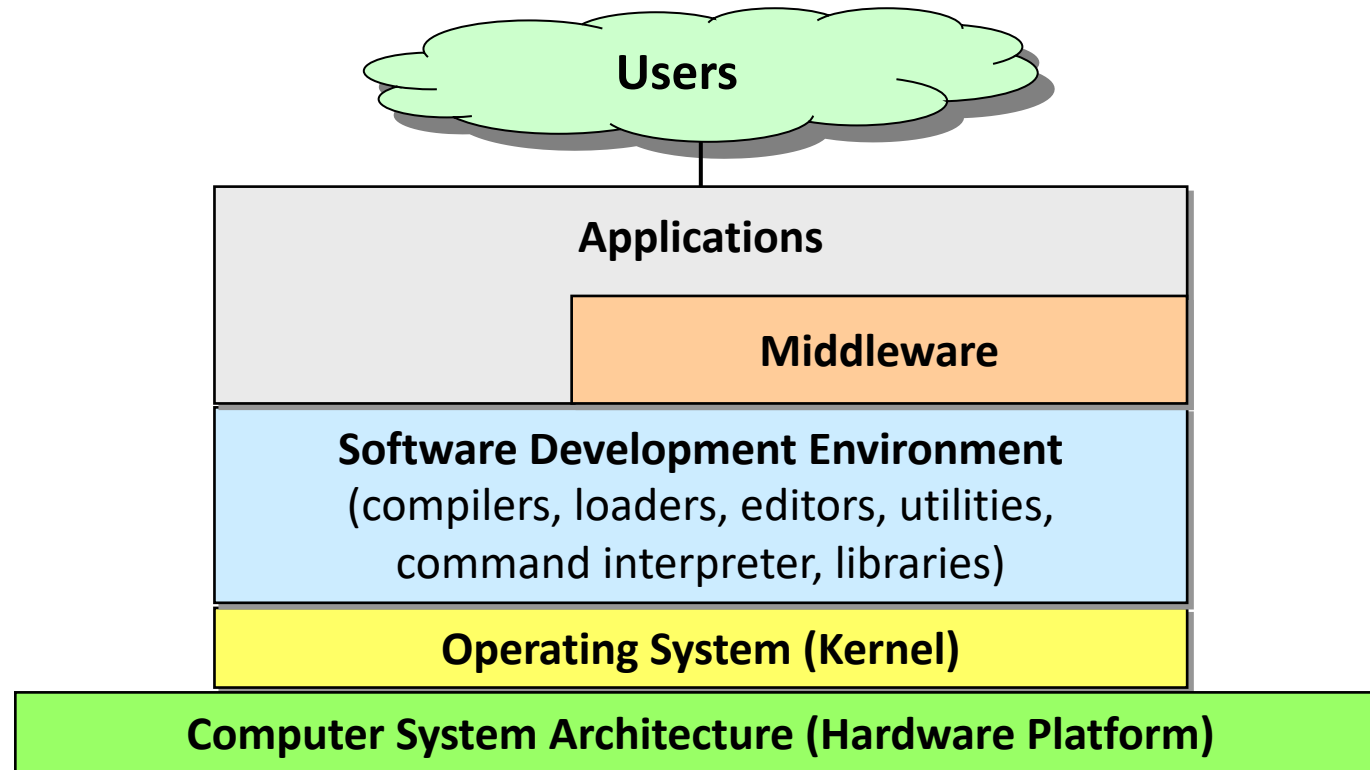




What is an Operating System?

Software Stack

- An OS is software that converts hardware into a useful form for applications



What happens when a program runs?

- A running program executes instructions
 - The processor fetches an instruction from memory
 - Decode: Figure out which instruction this is
 - Execute: i.e., add two numbers, access memory, check a condition, jump to function, and so forth
 - The processor moves on to the next instruction and so on
- Who did load this program into memory?
- Who does remove this program when this program acts abnormally?

What happens when many programs run together?

- Two programs share a single main memory
 - Program A may corrupt data of program B
 - How can I tell where my data are located at?
- Two programs share a small set of processors
 - When the number of cores is smaller than the number of running programs, who controls which program uses a processor?
 - Who decides which program runs on which core?
 - You have only one mouse and one keyboard, how can multiple programs share the resources?
- If a program tries to steal your information, which is stored by Chrome web browser, who protects your information from it?

Operating System (OS)

- Responsible for
 - Making it easy to **run** programs
 - Allowing programs to **share** memory
 - Enabling programs to **interact** with devices

OS is in charge of making sure the system operates
correctly and **efficiently**

Three Pieces



- Virtualization
 - Make each application believe it has each resource to itself
 - Processes, CPU scheduling, virtual memory
- Concurrency
 - Handle concurrent events correctly
 - Threads, synchronization
- Persistence
 - Access information permanently while preserving correctness upon unexpected failures
 - Storage, file systems

Virtualization

- The OS takes a physical resource and transforms it into a virtual form of itself
 - **Physical resource:** Processor, Memory, Disk ...
 - The virtual form is more general, powerful and easy-to-use
- An application cannot directly access (or see) the underlying hardware
 - An application sees and uses virtualized hardware
 - An OS is located in between an application and hardware, and acts like an intermediary

Roles of Operating System

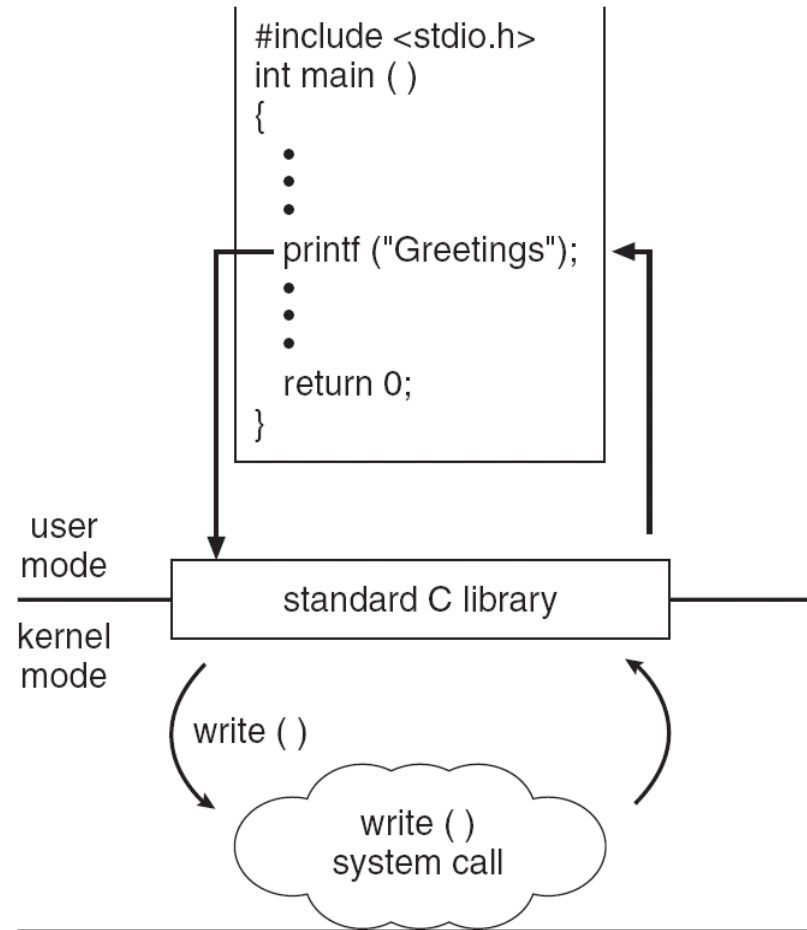


- OS is a **resource allocator**
 - Manages all resources
 - Decides between conflicting requests for efficient and fair resource use
- OS is a **control program**
 - Controls execution of programs to prevent errors and improper use of the computer

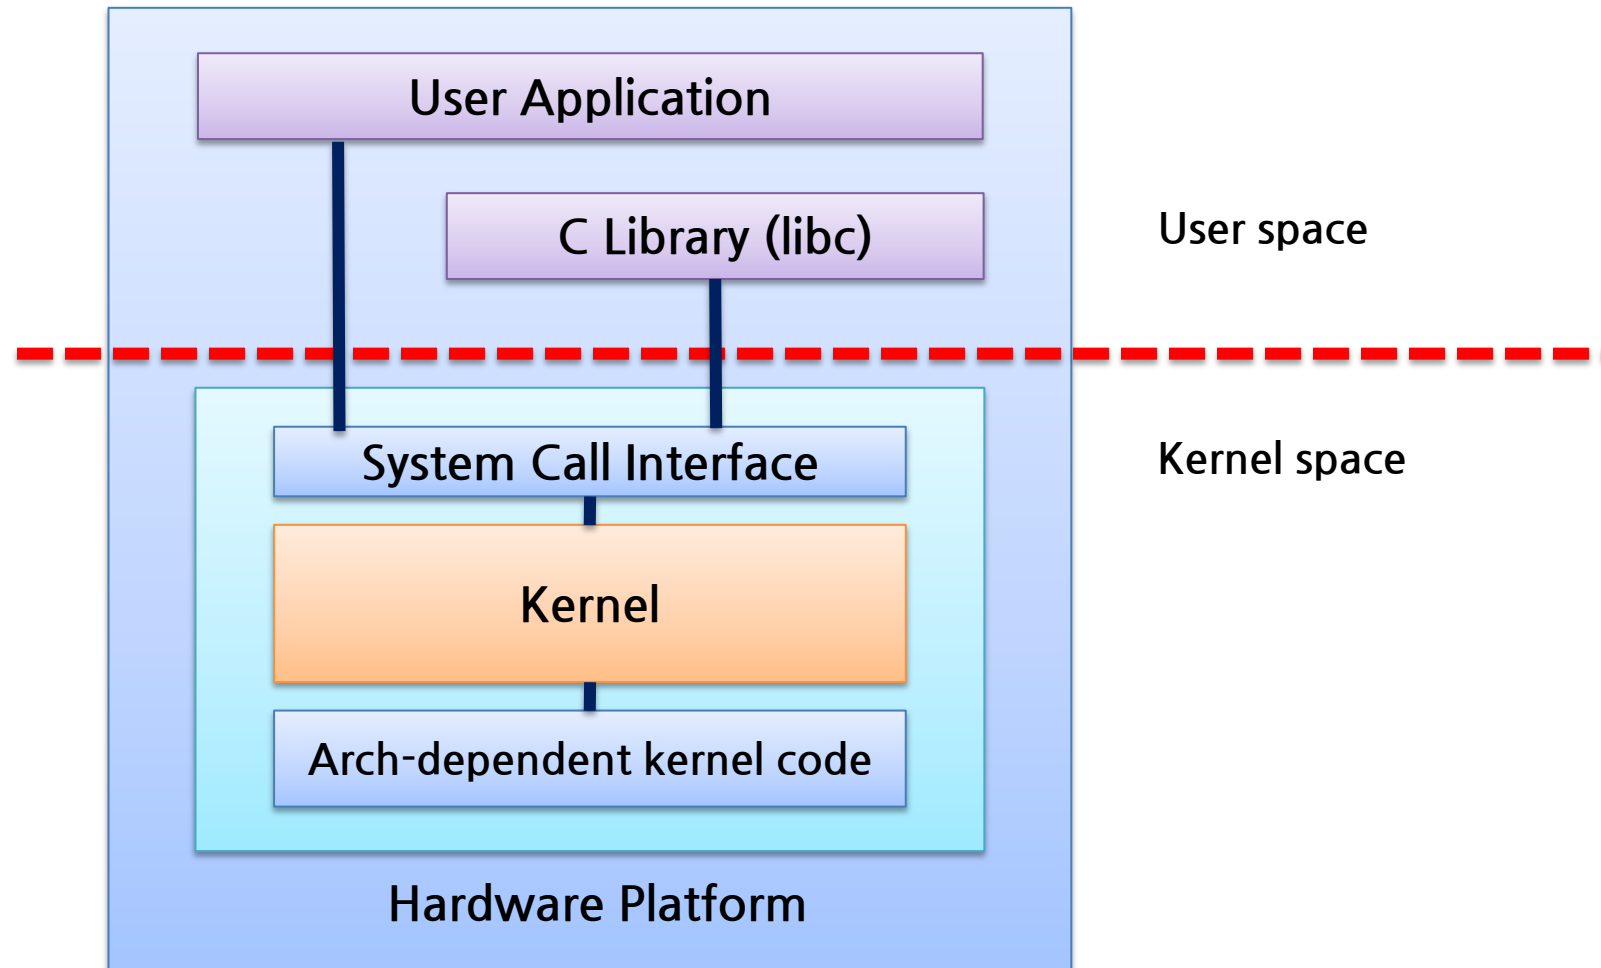
System calls

- A system call allows a program **to request OS what it wants to do**
 - Systems are interface between OS and applications
 - A typical OS exports a few hundred system calls.
 - Run programs
 - Access memory
 - Access devices

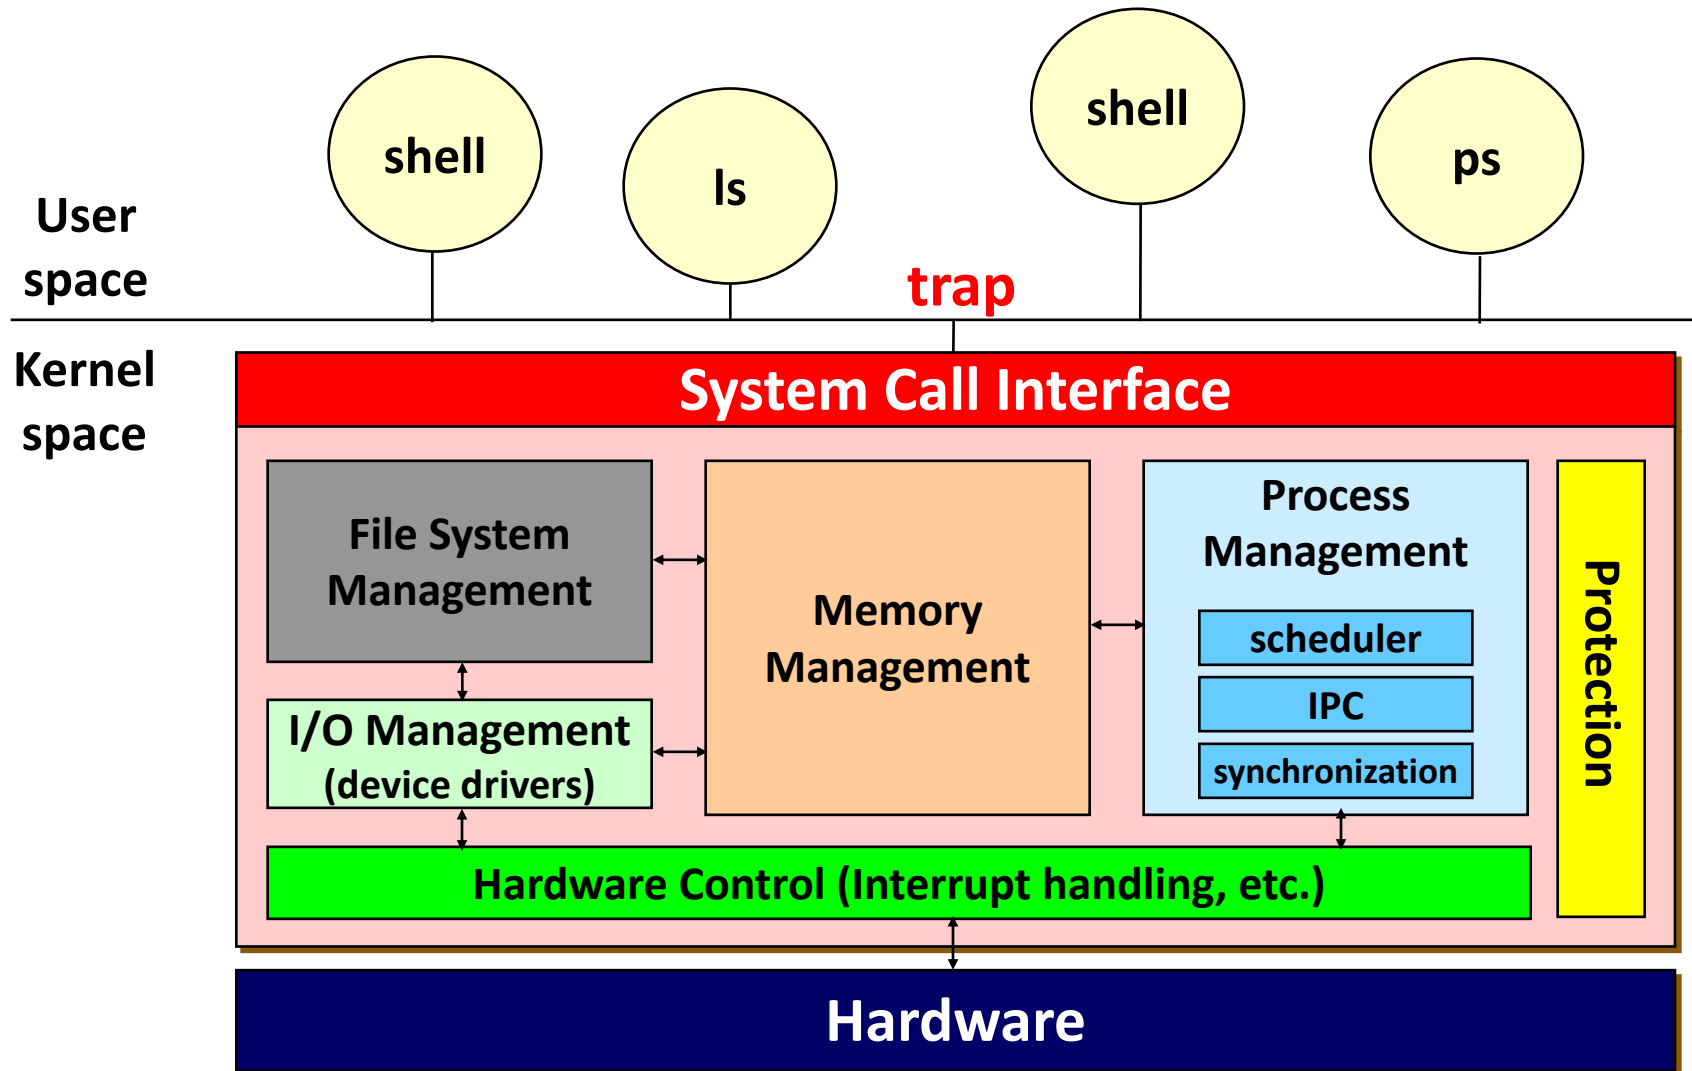
System Call Interface



Interactions Between OS and Apps



Interactions Between OS and Apps



Virtualizing CPU



- The system has a very large number of virtual CPUs
 - Turning a single CPU into a seemingly infinite number of CPUs
 - Allowing many programs to seemingly run at once
→ **Virtualizing the CPU**

Virtualizing CPU (Cont.)

```
1      #include <stdio.h>
2      #include <stdlib.h>
3      #include <sys/time.h>
4      #include <assert.h>
5      #include "common.h"
6
7      int
8      main(int argc, char *argv[])
9      {
10         if (argc != 2) {
11             fprintf(stderr, "usage: cpu <string>\n");
12             exit(1);
13         }
14         char *str = argv[1];
15         while (1) {
16             Spin(1); // Repeatedly checks the time and
17                     // returns once it has run for a second
18             printf("%s\n", str);
19         }
20         return 0;
21     }
```

Simple Example(cpu.c): Code That Loops and Prints

Virtualizing the CPU (Cont.)

- Execution result

```
prompt> gcc -o cpu cpu.c -Wall
prompt> ./cpu "A"
A
A
A
^C
prompt>
```

Run forever; Only by pressing "Control-c" can we halt the program

Virtualizing the CPU (Cont.)

- Execution result

```
prompt> ./cpu A & ; ./cpu B & ; ./cpu C & ; ./cpu D &  
[1] 7353  
[2] 7354  
[3] 7355  
[4] 7356  
A  
B  
D  
C  
A  
B  
D  
C  
A  
C  
B  
D  
...
```

Even though we have only **one processor**, all four of programs seem to be running **at the same time!**

Virtualizing Memory



- The physical memory is *an array of bytes*
- A program keeps all of its data structures in memory
 - **Read memory** (load):
 - Specify an address to be able to access the data
 - **Write memory** (store):
 - Specify the data to be written to the given address
- Let's read a program that accesses memory (`mem.c`)

Virtualizing Memory (Cont.)

```
1      #include <unistd.h>
2      #include <stdio.h>
3      #include <stdlib.h>
4      #include "common.h"
5
6      int
7      main(int argc, char *argv[])
8      {
9          int *p = malloc(sizeof(int)); // a1: allocate memory
10         assert(p != NULL);
11         printf("(d) address of p: %08x\n",
12                getpid(), (unsigned) p); // a2: print out the
                                           address of the memmory
13         *p = 0; // a3: put zero into the first slot of the memory
14         while (1) {
15             Spin(1);
16             *p = *p + 1;
17             printf("(d) p: %d\n", getpid(), *p); // a4
18         }
19         return 0;
20     }
```

Virtualizing Memory (Cont.)



- The output of the program `mem.c`

```
prompt> ./mem
(2134) memory address of p: 00200000
(2134) p: 1
(2134) p: 2
(2134) p: 3
(2134) p: 4
(2134) p: 5
^C
```

- The newly allocated memory is at address `00200000`.
- It updates the value and prints out the result.

Virtualizing Memory (Cont.)

- Running `mem.c` multiple times

```
prompt> ./mem & ./mem &  
[1] 24113  
[2] 24114  
(24113) memory address of p: 00200000  
(24114) memory address of p: 00200000  
(24113) p: 1  
(24114) p: 1  
(24114) p: 2  
(24113) p: 2  
(24113) p: 3  
(24114) p: 3  
...
```

- It is as if each running program has its **own private memory**.
 - Each running program has allocated memory at the same address.
 - Each seems to be updating the value at `00200000` independently.

Virtualizing Memory (Cont.)



- Each process accesses its own private **virtual address space**
 - The OS maps **address space** onto the **physical memory**
 - A memory reference within one running program does not affect the address space of other processes
 - Physical memory is a shared resource, managed by the OS

Problem of Concurrency

- The OS is juggling **many things at once**, first running one process, then another, and so forth
- Modern **multi-threaded programs** also exhibit the concurrency problem

Concurrency Example

■ A Multi-threaded Program (thread.c)

```
1      #include <stdio.h>
2      #include <stdlib.h>
3      #include "common.h"
4
5      volatile int counter = 0;
6      int loops;
7
8      void *worker(void *arg) {
9          int i;
10         for (i = 0; i < loops; i++) {
11             counter++;
12         }
13         return NULL;
14     }
15
16     int
17     main(int argc, char *argv[])
18     {
19         if (argc != 2) {
20             fprintf(stderr, "usage: threads <value>\n");
21             exit(1);
22         }
```

Concurrency Example (Cont.)

```
16     int
17     main(int argc, char *argv[])
18     {
19         if (argc != 2) {
20             fprintf(stderr, "usage: threads <value>\n");
21             exit(1);
22         }
23         loops = atoi(argv[1]);
24         pthread_t p1, p2;
25         printf("Initial value : %d\n", counter);
26
27         Pthread_create(&p1, NULL, worker, NULL);
28         Pthread_create(&p2, NULL, worker, NULL);
29         Pthread_join(p1, NULL);
30         Pthread_join(p2, NULL);
31         printf("Final value : %d\n", counter);
32         return 0;
33     }
```

- The main program creates **two threads**.
 - Thread: a function running within the same memory space. Each thread start running in a routine called `worker()`
 - `worker()`: increments a counter

Concurrency Example (Cont.)

- `loops` determines how many times each of the two workers will **increment the shared counter** in a loop

- `loops: 1000`

```
prompt> gcc -o thread thread.c -Wall -pthread
prompt> ./thread 1000
Initial value : 0
Final value : 2000
```

- `loops: 100000`

```
prompt> ./thread 100000
Initial value : 0
Final value : 143012 // huh??
prompt> ./thread 100000
Initial value : 0
Final value : 137298 // what the??
```

Why is this happening?

- Increment a shared counter → take three instructions.
 1. Load the value of the counter from memory into register.
 2. Increment it
 3. Store it back into memory
- These three instructions do not execute **atomically** → Problem of **concurrency** happen

Persistence

- Devices such as DRAM store values in a volatile
- *Hardware* and *software* are needed to store data **persistently**
 - **Hardware:** I/O device such as a hard drive, solid-state drives(SSDs)
 - **Software:**
 - File system manages the disk
 - File system is responsible for storing any files the user creates

Persistence (Cont.)

- Create a file (`/tmp/file`) that contains the string "hello world"

```
1      #include <stdio.h>
2      #include <unistd.h>
3      #include <assert.h>
4      #include <fcntl.h>
5      #include <sys/types.h>
6
7      int
8      main(int argc, char *argv[])
9      {
10         int fd = open("/tmp/file", O_WRONLY | O_CREAT
                        | O_TRUNC, S_IRWXU);
11         assert(fd > -1);
12         int rc = write(fd, "hello world\n", 13);
13         assert(rc == 13);
14         close(fd);
15         return 0;
16     }
```

`open()`, `write()`, and `close()` system calls are routed to the part of OS and called the file system, which handles the requests

Persistence (Cont.)

- A disk, like main memory, is just a large array of bytes
- What OS does in order to write data to disk and manage them?
 - Figure out **where** on disk this new data will reside
 - **Issue I/O** requests to the underlying storage device
- File system handles system crashes during write
 - **Journaling** or **copy-on-write**
 - Carefully ordering writes to disk

Design Goals

- Build up **abstraction**
 - Make the system convenient and easy to use
- Provide high **performance**
 - Minimize the overhead of the OS
 - OS must strive to provide virtualization without excessive overhead
- **Protection** between applications
 - Isolation: Bad behavior of one does not harm other and the OS itself

Design Goals (Cont.)



- High degree of **reliability**
 - The OS must also run non-stop
- Other issues
 - Energy-efficiency
 - Security
 - Mobility