



Program and Process

The fork() System Call

- Create a new process
 - The newly-created process has its own copy of the **address space, registers, and PC**

p1.c

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char *argv[]) {
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {          // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) { // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else {              // parent goes down this path (main)
        printf("hello, I am parent of %d (pid:%d)\n",
            rc, (int) getpid());
    }
    return 0;
}
```

Calling fork() example (Cont.)

Result (Not deterministic)

```
prompt> ./p1
hello world (pid:29146)
hello, I am parent of 29147 (pid:29146)
hello, I am child (pid:29147)
prompt>
```

or

```
prompt> ./p1
hello world (pid:29146)
hello, I am child (pid:29147)
hello, I am parent of 29147 (pid:29146)
prompt>
```

The wait() System Call

- This system call won't return until the child has run and exited

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main(int argc, char *argv[]) {
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {                // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {        // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else {                    // parent goes down this path (main)
        int wc = wait(NULL);
        printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
            rc, wc, (int) getpid());
    }
    return 0;
}
```

p2.c

The wait() System Call (Cont.)



Result (Deterministic)

```
prompt> ./p2
hello world (pid:29266)
hello, I am child (pid:29267)
hello, I am parent of 29267 (wc:29267) (pid:29266)
prompt>
```

The exec() System Call

- Run a program that is different from the parent

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <sys/wait.h>

int main(int argc, char *argv[]) {
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
        char *myargs[3];
        myargs[0] = strdup("wc");
        myargs[1] = strdup("p3.c");
        myargs[2] = NULL;
        ...
    }
}
```

p3.c

// program: "wc" (word count)
// argument: file to count
// marks end of array

The exec() System Call (Cont.)

p3.c (Cont.)

```
...
    execvp(myargs[0], myargs); // runs word count
    printf("this shouldn't print out");
} else { // parent goes down this path (main)
    int wc = wait(NULL);
    printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
          rc, wc, (int) getpid());
}
return 0;
}
```

Result

```
prompt> ./p3
hello world (pid:29383)
hello, I am child (pid:29384)
29 107 1030 p3.c
hello, I am parent of 29384 (wc:29384) (pid:29383)
prompt>
```

How to provide the illusion of many CPUs?

- CPU virtualizing
 - The OS can promote the illusion that many virtual CPUs exist
 - **Time sharing**: Running one process, then stopping it and running another
 - The potential cost is **performance**

A Process


A process is a **program in execution**

- Comprising of a process:
 - Memory (address space)
 - Instructions
 - Data section
 - Registers
 - Program counter
 - Stack pointer
- A process is different than a program
 - Program: Static code and static data
 - Process: Dynamic instance of code and data
- Can have multiple process instances of same program
 - Can have multiple processes of the same program
 - Example: many users can run "ls" at the same time

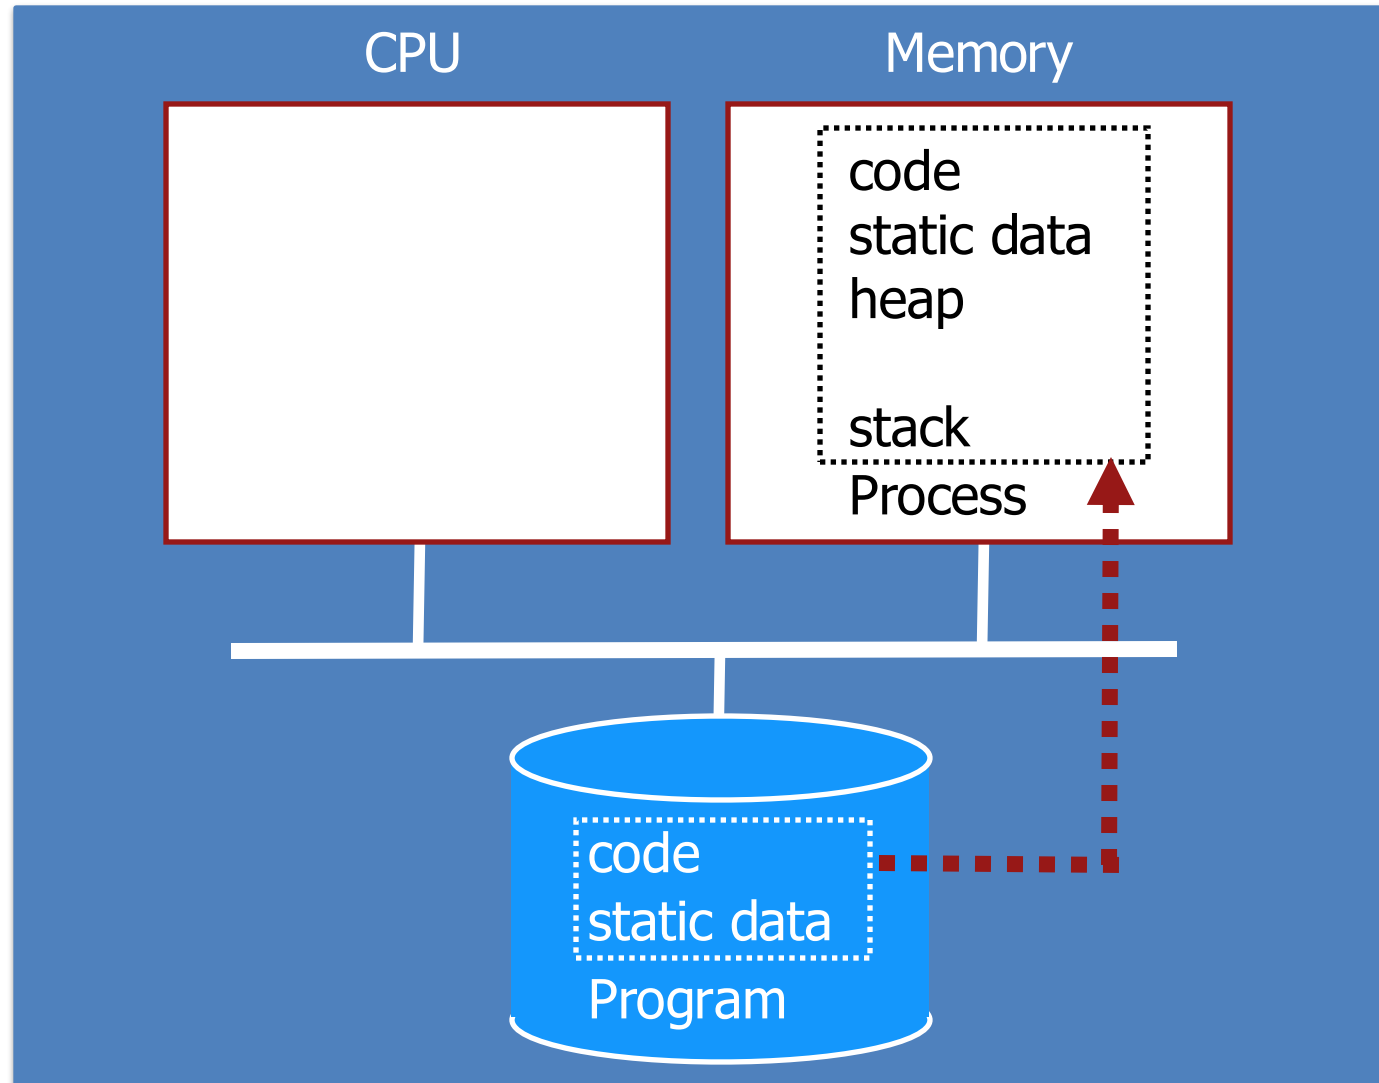
Process API

- These APIs are available on any modern OS.
 - **Create**
 - Create a new process to run a program
 - **Destroy**
 - Halt a runaway process
 - **Wait**
 - Wait for a process to stop running
 - **Miscellaneous Control**
 - Some kind of method to suspend a process and then resume it
 - **Status**
 - Get some status info about a process

Process Creation



Process Creation



Process Creation

1. **Load** a program code into memory, into the address space of the process
 - Programs initially reside on disk in *executable format*
 - OS perform the loading process **lazily**
 - Loading pieces of code or data only as they are needed during program execution
2. The program's run-time **stack** is allocated
 - Use the stack for *local variables, function parameters, and return address*
 - Initialize the stack with arguments → `argc` and the `argv` array of `main()` function

Process Creation (Cont.)

3. The program's **heap** is created

- Used for explicitly requested dynamically allocated data
- Program requests such space by calling `malloc()` and frees it by calling `free()`


4. The OS does some other initialization tasks

- input/output (I/O) setup
 - Each process by default has three open file descriptors
 - Standard input, output and error

5. **Start the program** running at the entry point, namely `main()`

- The OS *transfers control* of the CPU to the newly-created process

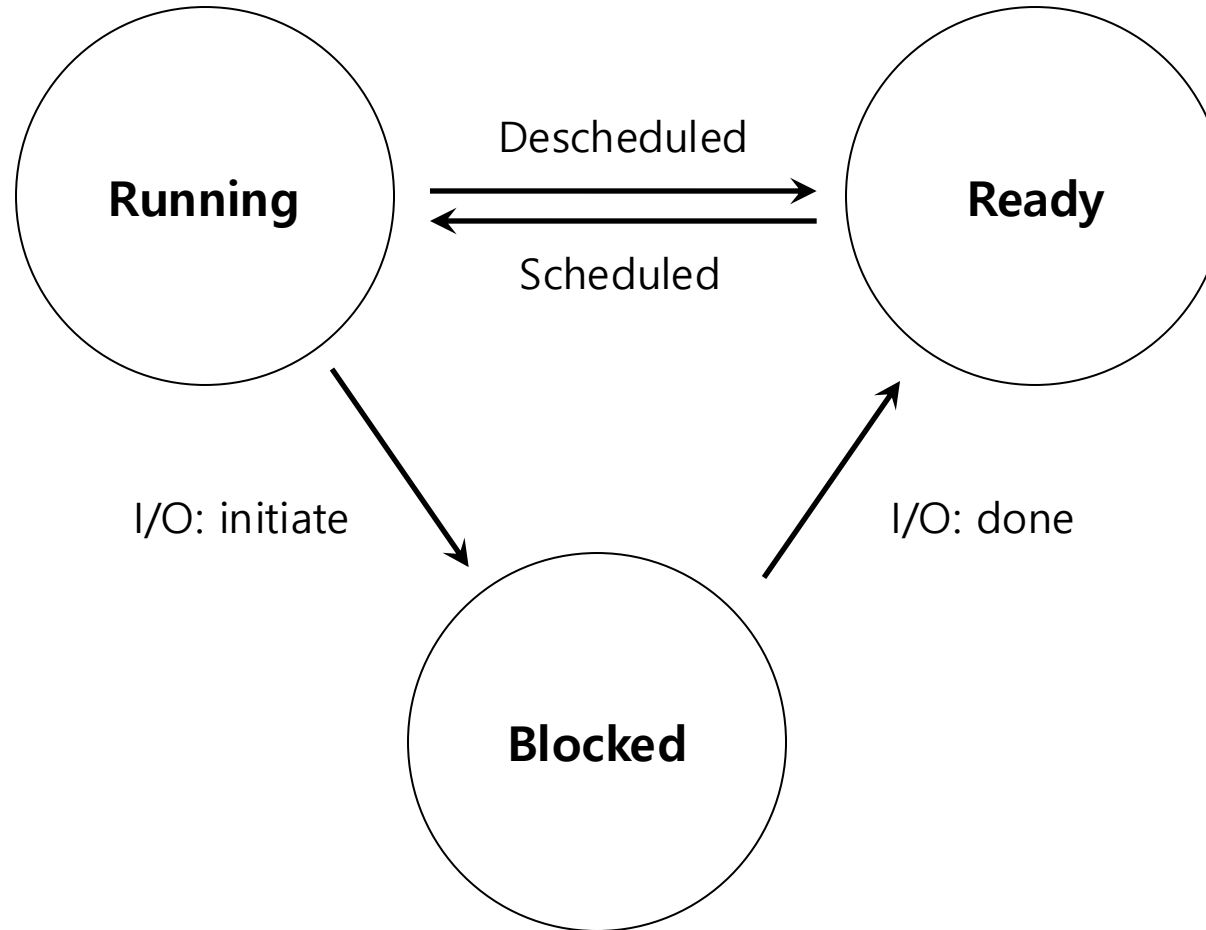
Loading: From Program To Process



Process States

- A process can be one of three states
 - **Running**
 - A process is running on a processor
 - **Ready**
 - A process is ready to run but for some reason the OS has chosen not to run it at this given moment
 - **Blocked**
 - A process has performed some kind of operation
 - When a process initiates an I/O request to a disk, it becomes blocked and thus some other process can use the processor

Process State Transition



Data structures

- The OS has some key data structures that track various relevant pieces of information
 - **Process list**
 - Ready processes
 - Blocked processes
 - Current running process
 - **Register context**
- PCB(Process Control Block)
 - A C-structure that contains information about each process

Example) The xv6 kernel Proc Structure

```
// the registers xv6 will save and restore
// to stop and subsequently restart a process
struct context {
    int eip;    // Index pointer register
    int esp;    // Stack pointer register
    int ebx;    // Called the base register
    int ecx;    // Called the counter register
    int edx;    // Called the data register
    int esi;    // Source index register
    int edi;    // Destination index register
    int ebp;    // Stack base pointer register
};

// the different states a process can be in
enum proc_state { UNUSED, EMBRYO, SLEEPING,
                  RUNNABLE, RUNNING, ZOMBIE };
```

Example) The xv6 kernel Proc Structure

```
// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char *mem;                // Start of process memory
    uint sz;                  // Size of process memory
    char *kstack;             // Bottom of kernel stack
                                // for this process
    enum proc_state state;    // Process state
    int pid;                  // Process ID
    struct proc *parent;      // Parent process
    void *chan;               // If non-zero, sleeping on chan
    int killed;               // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;         // Current directory
    struct context context;    // Switch here to run process
    struct trapframe *tf;      // Trap frame for the
                                // current interrupt
};
```

How to efficiently virtualize the CPU with control?

- The OS needs to share the physical CPU by time sharing
- Issue
 - Performance: How can we implement virtualization without adding excessive overhead to the system?
 - Control: How can we run processes efficiently while retaining control over the CPU?

Direct Execution

- ▣ Just run the program directly on the CPU

OS	Program
<ol style="list-style-type: none">1. Create entry for process list2. Allocate memory for program3. Load program into memory4. Set up stack with <code>argc / argv</code>5. Clear registers6. Execute <code>call main()</code>	<ol style="list-style-type: none">7. Run <code>main()</code>8. Execute <code>return from main()</code>
<ol style="list-style-type: none">9. Free memory of process10. Remove from process list	

**Without *limits* on running programs,
the OS wouldn't be in control of anything and
thus would be "just a library"**

Problem 1: Restricted Operation

- What if a process wishes to perform some kind of restricted operation such as ...
 - Issuing an I/O request to a disk
 - Gaining access to more system resources such as CPU or memory
- **Solution:** Using protected control transfer
 - **User mode:** Applications do not have full access to hardware resources
 - **Kernel mode:** The OS has access to the full resources of the machine

System Call

- Allow the kernel to **carefully expose** certain key pieces of functionality to user program, such as ...
 - Accessing the file system
 - Creating and destroying processes
 - Communicating with other processes
 - Allocating more memory

System Call (Cont.)

- **Trap** instruction
 - Jump into the kernel
 - Raise the privilege level to kernel mode
- **Return-from-trap** instruction
 - Return into the calling user program
 - Reduce the privilege level back to user mode

Limited Direction Execution Protocol



OS @ boot
(kernel mode)

Hardware

initialize trap table

remember address of ...
syscall handler

OS @ run
(kernel mode)

Hardware

Program
(user mode)

Create entry for process list
Allocate memory for program
Load program into memory
Setup user stack with argv
Fill kernel stack with reg/PC
return-from-trap

restore regs from kernel stack
move to user mode
jump to main

Run main()
...
Call system
trap into OS

Limited Direction Execution Protocol



OS @ run
(kernel mode)

Hardware

Program
(user mode)

(Cont.)

Handle trap
Do work of syscall
return-from-trap

save regs to kernel stack
move to kernel mode
jump to trap handler

restore regs from kernel stack
move to user mode
jump to PC after trap

...
return from main
trap (via `exit()`)

Free memory of process
Remove from process list

Problem 2: Switching Between Processes

- How can the OS regain control of the CPU so that it can switch between processes?
 - A cooperative Approach: Wait for system calls
 - A Non-Cooperative Approach: The OS takes control

A cooperative Approach: Wait for system calls

- Processes **periodically give up the CPU** by making **system calls** such as `yield`
 - The OS decides to run some other task
 - Application also transfer control to the OS when they do something illegal
 - Divide by zero
 - Try to access memory that it shouldn't be able to access
 - Ex) Early versions of the Macintosh OS, The old Xerox Alto system

**A process gets stuck in an infinite loop
→ Reboot the machine**

A Non-Cooperative Approach: OS Takes Control

- A timer interrupt
 - During the boot sequence, the OS start the timer.
 - The timer raise an *interrupt* every so many milliseconds.
 - When the interrupt is raised :
 - The currently running process is halted.
 - Save enough of the state of the program
 - A pre-configured interrupt handler in the OS runs.

A timer interrupt gives OS the ability to run again on a CPU.

Saving and Restoring Context

- Scheduler makes a decision:
 - Whether to continue running the **current process**, or switch to a **different one**.
 - If the decision is made to switch, the OS executes context switch

Context Switch



- A low-level piece of assembly code
 - **Save a few register values** for the current process onto its kernel stack
 - General purpose registers
 - PC
 - kernel stack pointer
 - **Restore a few** for the soon-to-be-executing process from its kernel stack
 - **Switch to the kernel stack** for the soon-to-be-executing process

Limited Direction Execution Protocol (Timer interrupt)

OS @ boot (kernel mode)	Hardware	
initialize trap table	remember address of ... syscall handler timer handler	
start interrupt timer	start timer interrupt CPU in X ms	
OS @ run (kernel mode)	Hardware	Program (user mode)
		Process A ...
	timer interrupt save regs(A) to k-stack(A) move to kernel mode jump to trap handler	

Limited Direction Execution Protocol (Timer interrupt)

OS @ run
(kernel mode)

Hardware

Program
(user mode)

(Cont.)

Handle the trap
Call switch() routine
 save regs(A) to proc-struct(A)
 restore regs(B) from proc-struct(B)
 switch to k-stack(B)
return-from-trap (into B)

restore regs(B) from k-stack(B)
move to user mode
jump to B's PC

Process B

...

The xv6 Context Switch Code

```
1 # void swtch(struct context *old, struct context *new);
2 #
3 # Save current register context in old
4 # and then load register context from new.
5 .globl swtch
6 swtch:
7     # Save old registers
8     movl 4(%esp), %eax           # put old ptr into eax
9     popl 0(%eax)                # save the old IP
10    movl %esp, 4(%eax)           # and stack
11    movl %ebx, 8(%eax)           # and other registers
12    movl %ecx, 12(%eax)
13    movl %edx, 16(%eax)
14    movl %esi, 20(%eax)
15    movl %edi, 24(%eax)
16    movl %ebp, 28(%eax)
17
18    # Load new registers
19    movl 4(%esp), %eax           # put new ptr into eax
20    movl 28(%eax), %ebp          # restore other registers
21    movl 24(%eax), %edi
22    movl 20(%eax), %esi
23    movl 16(%eax), %edx
24    movl 12(%eax), %ecx
25    movl 8(%eax), %ebx
26    movl 4(%eax), %esp          # stack is switched here
27    pushl 0(%eax)               # return addr put in place
28    ret                         # finally return into new ctxt
```

Worried About Concurrency?

- What happens if, during interrupt or trap handling, another interrupt occurs?
- OS handles these situations:
 - **Disable interrupts** during interrupt processing
 - Use a number of sophisticated **locking** schemes to protect concurrent access to internal data structures